

# Data Transformations: ncyflights13\$flights\_1

Monica Buczynski

08/05/2021

Note: The purpose of this document is to showcase a sample of skills covered in *R for Data Science* (chapter: Data Transformations) by Garrett Golemund and Hadley Wickham. All scripts were taken from <https://r4ds.had.co.nz/transform.html> and <https://jrnold.github.io/r4ds-exercise-solutions/index.html>. The code for each exercise was studied carefully for understanding and then was retyped manually into R to maximize the learning experience; however, many of the original scripts were altered for further experimentation and presentation aesthetics.

The skills that I focused on include:

- Filter rows with *filter()*
- Arrange rows with *arrange()*
- Select columns with *select()*
- Add new variables with *mutate()*
- Grouped summaries with *summarise()*
- Grouped mutates (and filters)

```
# View first row of data
head(flights, 1)
```

```
## # A tibble: 1 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>    <int>         <int>
## 1  2013     1     1     517             515         2      830             819
## # ... with 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

Find flights that arrived more than two hours late, but didn't leave late.

```
head(filter(flights, dep_delay <= 0, arr_delay > 120))
```

```
## # A tibble: 6 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>    <int>         <int>
## 1  2013     1    27    1419             1420        -1     1754             1550
## 2  2013    10     7    1350             1350         0     1736             1526
## 3  2013    10     7    1357             1359        -2     1858             1654
## 4  2013    10    16     657             700        -3     1258             1056
## 5  2013    11     1     658             700        -2     1329             1015
## 6  2013     3    18    1844             1847        -3         39             2219
## # ... with 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

Find flights that flew to Houston (IAH or HOU).

```
head(filter(flights, dest == "IAH" | dest == "HOU")) # characters need quotation marks
```

```
## # A tibble: 6 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>    <int>         <int>
## 1  2013     1     1     517             515         2      830             819
## 2  2013     1     1     533             529         4      850             830
## 3  2013     1     1     623             627        -4      933             932
## 4  2013     1     1     728             732        -4     1041             1038
## 5  2013     1     1     739             739         0     1104             1038
## 6  2013     1     1     908             908         0     1228             1219
## # ... with 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

Find flights that were operated by United, American, or Delta.

```
airlines # to lookup airline codes
```

```
## # A tibble: 16 x 2
##   carrier name
##   <chr>   <chr>
## 1 9E      Endeavor Air Inc.
## 2 AA      American Airlines Inc.
## 3 AS      Alaska Airlines Inc.
## 4 B6      JetBlue Airways
## 5 DL      Delta Air Lines Inc.
## 6 EV      ExpressJet Airlines Inc.
## 7 F9      Frontier Airlines Inc.
## 8 FL      AirTran Airways Corporation
## 9 HA      Hawaiian Airlines Inc.
## 10 MQ     Envoy Air
## 11 OO     SkyWest Airlines Inc.
## 12 UA     United Air Lines Inc.
## 13 US     US Airways Inc.
## 14 VX     Virgin America
## 15 WN     Southwest Airlines Co.
## 16 YV     Mesa Airlines Inc.
```

```
head(filter(flights, carrier == "UA" | carrier == "AA" | carrier == "DL"))
```

```
## # A tibble: 6 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>       <dbl>   <int>         <int>
## 1  2013     1     1     517           515         2     830           819
## 2  2013     1     1     533           529         4     850           830
## 3  2013     1     1     542           540         2     923           850
## 4  2013     1     1     554           600        -6     812           837
## 5  2013     1     1     554           558        -4     740           728
## 6  2013     1     1     558           600        -2     753           745
## # ... with 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

Find flights that do NOT have a greater delay than 120 minutes.

```
head(filter(flights, !(arr_delay > 120 | dep_delay > 120)))
```

```
## # A tibble: 6 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>       <dbl>   <int>         <int>
## 1  2013     1     1     517           515         2     830           819
## 2  2013     1     1     533           529         4     850           830
## 3  2013     1     1     542           540         2     923           850
## 4  2013     1     1     544           545        -1    1004          1022
## 5  2013     1     1     554           600        -6     812           837
## 6  2013     1     1     554           558        -4     740           728
## # ... with 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

Find flights that do have a delay that is less than 120 minutes.

```
head(filter(flights, arr_delay <= 120, dep_delay <= 120))
```

```
## # A tibble: 6 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>       <dbl>   <int>         <int>
## 1  2013     1     1     517           515         2     830           819
## 2  2013     1     1     533           529         4     850           830
## 3  2013     1     1     542           540         2     923           850
## 4  2013     1     1     544           545        -1    1004          1022
## 5  2013     1     1     554           600        -6     812           837
## 6  2013     1     1     554           558        -4     740           728
## # ... with 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

Find flights that had an arrival delay of two or more hours.

```
head(filter(flights, arr_delay >=120))
```

```
## # A tibble: 6 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>       <dbl>   <int>         <int>
## 1  2013     1     1     811           630       101    1047           830
## 2  2013     1     1     848          1835       853    1001          1950
## 3  2013     1     1     957           733       144    1056           853
## 4  2013     1     1    1114           900       134    1447          1222
## 5  2013     1     1    1505          1310       115    1638          1431
## 6  2013     1     1    1525          1340       105    1831          1626
## # ... with 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

Find flights that were delayed by at least an hour, but made up over 30 minutes in flight.

```
head(filter(flights, dep_delay >= 60, dep_delay - arr_delay > 30))
```

```
## # A tibble: 6 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>       <dbl>   <int>         <int>
## 1  2013     1     1    2205          1720       285     46          2040
## 2  2013     1     1    2326          2130       116    131           18
## 3  2013     1     3    1503          1221       162    1803          1555
## 4  2013     1     3    1839          1700        99    2056          1950
## 5  2013     1     3    1850          1745        65    2148          2120
## 6  2013     1     3    1941          1759       102    2246          2139
## # ... with 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

Find flights that departed between midnight and 6am (inclusive).

```
summary(flights$dep_time) # to find if midnight is denoted as 2400 or 0.
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.     NA's
##         1      907     1401    1349    1744    2400     8255
```

```
head(filter(flights, dep_time == 2400 | dep_time <= 6000))
```

```
## # A tibble: 6 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>    <int>         <int>
## 1  2013     1     1     517           515         2      830           819
## 2  2013     1     1     533           529         4      850           830
## 3  2013     1     1     542           540         2      923           850
## 4  2013     1     1     544           545        -1     1004          1022
## 5  2013     1     1     554           600        -6      812           837
## 6  2013     1     1     554           558        -4      740           728
## # ... with 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

Sort flights to find the most delayed flight in the dataset.

Flight HA 5 from JFK to HNL had a 1301 minute (21.68 hours) delay on January 9th, 2013.

```
slice_head(flights %>%
  select(dep_delay, carrier, flight, origin, dest, month, day, year) %>%
  arrange(desc(dep_delay)))
```

```
## # A tibble: 1 x 8
##   dep_delay carrier flight origin dest  month   day  year
##   <dbl> <chr>    <int> <chr> <chr> <int> <int> <int>
## 1    1301 HA         51 JFK   HNL    1     9  2013
```

Find the flight that left the earliest in the dataset.

Flight B6 97 from JFK to DEN departed 43 minutes early than scheduled on December 7th, 2013.

```
slice_head(flights %>%
  select(dep_delay, carrier, flight, origin, dest, month, day, year) %>%
  arrange(dep_delay))
```

```
## # A tibble: 1 x 8
##   dep_delay carrier flight origin dest  month   day  year
##   <dbl> <chr>    <int> <chr> <chr> <int> <int> <int>
## 1     -43 B6         97 JFK   DEN   12     7  2013
```

Sort flights to find the fastest (highest speed) flight in the dataset.

Flight DL 1499 had the fastest average ground speed of 703.38 miles/hour.

```
slice_head(flights %>%
  mutate(ground_speed = distance/(air_time/60)) %>% # create a new variable, ground_speed
  arrange(desc(distance/air_time)) %>%
  select(ground_speed, carrier, flight, origin, dest, month, day, year))
```

```
## # A tibble: 1 x 8
##   ground_speed carrier flight origin dest month day year
##   <dbl> <chr>    <int> <chr> <chr> <int> <int> <int>
## 1      703. DL      1499 LGA   ATL     5    25  2013
```

Which flights traveled the farthest?

Flight HA 51 from JFK to HNL is the longest #flight with a distance of 4,983 miles.

```
slice_head((flights %>%
  select(distance, carrier, flight, origin, dest) %>%
  arrange(desc(distance))))
```

```
## # A tibble: 1 x 5
##   distance carrier flight origin dest
##   <dbl> <chr>    <int> <chr> <chr>
## 1   4983 HA      51 JFK   HNL
```

Which flight traveled the shortest?

Flight US 1632 from EWR to LGA is the shortest flight with a distance of 17 miles.

```
slice_head(flights %>%
  select(distance, carrier, flight, origin, dest) %>%
  arrange(distance))
```

```
## # A tibble: 1 x 5
##   distance carrier flight origin dest
##   <dbl> <chr>    <int> <chr> <chr>
## 1     17 US      1632 EWR   LGA
```

For the flights that have a missing `dep_time`, what other variables are missing? What might these other missing rows represent?

Since `arrive_time` is also missing, these may be canceled flights.

```
head(filter(flights, is.na(dep_time)))
```

```
## # A tibble: 6 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     1     NA             1630           NA       NA             1815
## 2  2013     1     1     NA             1935           NA       NA             2240
## 3  2013     1     1     NA             1500           NA       NA             1825
## 4  2013     1     1     NA              600           NA       NA              901
## 5  2013     1     2     NA             1540           NA       NA             1747
## 6  2013     1     2     NA             1620           NA       NA             1746
## # ... with 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

Identify flights which were not canceled.

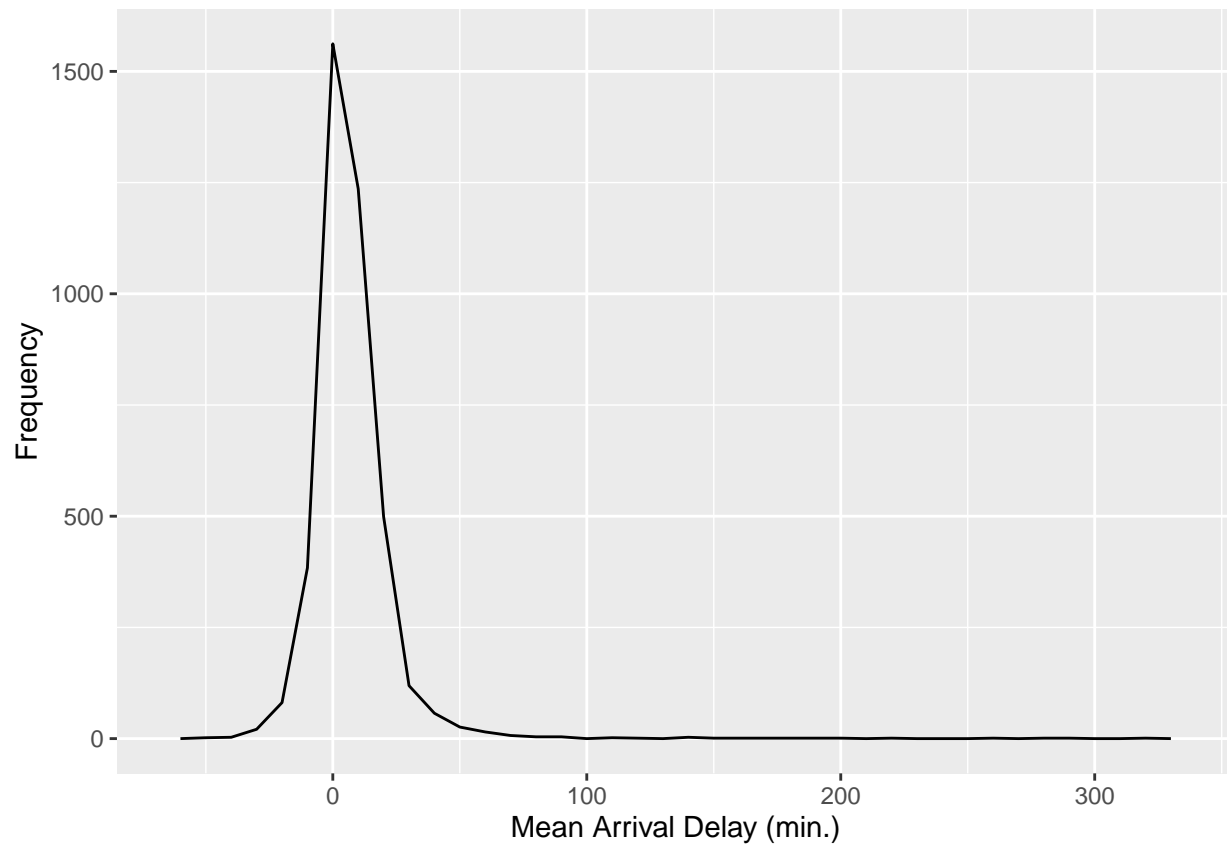
```
not_cancelled <- filter(flights, !is.na(dep_time)|!is.na(arr_time))
```

Planes (identified by their tail number) that have the highest average delays

```
head(delays <- not_cancelled %>%  
  filter(!is.na(arr_delay)) %>%  
  group_by(tailnum) %>%  
  summarise(mean_arr_delay = mean(arr_delay)) %>%  
  arrange(desc(mean_arr_delay)), 10)
```

```
## # A tibble: 10 x 2  
##   tailnum mean_arr_delay  
##   <chr>         <dbl>  
## 1 N844MH         320  
## 2 N911DA         294  
## 3 N922EV         276  
## 4 N587NW         264  
## 5 N851NW         219  
## 6 N928DN         201  
## 7 N7715E         188  
## 8 N654UA         185  
## 9 N665MQ         175.  
## 10 N427SW         157
```

```
ggplot(data = delays, aes(x= mean_arr_delay)) +  
  geom_freqpoly(binwidth = 10) +  
  xlab("Mean Arrival Delay (min.)") + ylab("Frequency")
```






*#scatterplot of number of flights vs. average delay*

```
head(delays <- not_cancelled %>%  
  filter(!is.na(arr_delay)) %>%  
  group_by(tailnum) %>%  
  summarise(delay = mean(arr_delay), n=n()))
```


```
## # A tibble: 6 x 3  
##   tailnum delay    n  
##   <chr>   <dbl> <int>  
## 1 D942DN  31.5     4  
## 2 NOEGMQ   9.98   352  
## 3 N10156  12.7   145  
## 4 N102UW   2.94    48  
## 5 N103US  -6.93    46  
## 6 N104UW   1.80    46
```

```
ggplot(data = delays, mapping = aes(x=n, y =delay))+ geom_point(alpha=1/10)
```



```
# useful to filter out the groups with the smallest numbers of observations,  
#so you can see more of the pattern  
#and less of the extreme variation in the smallest groups.
```

```
delays %>%  
  filter(n>25) %>%  
  ggplot(mapping = aes(x=n, y=delay)) +  
  geom_point(alpha=1/10)
```



Using measures of location: mean(x), median(x)

```
head(not_cancelled %>%
  filter(!is.na(arr_delay)) %>%
  group_by(tailnum) %>%
  summarise(mean_arr_delay = mean(arr_delay[arr_delay > 0]),
            mean_arr_early = mean(arr_delay[arr_delay < 0])) %>%
  select(tailnum, mean_arr_delay, mean_arr_early, everything()) %>%
  arrange(desc(mean_arr_delay))
```

```
## # A tibble: 6 x 3
##   tailnum mean_arr_delay mean_arr_early
##   <chr>         <dbl>         <dbl>
## 1 N104UW         45.7           -15.5
## 2 D942DN         45.7            -11
## 3 N102UW         40.6           -12.6
## 4 N10156         39.3           -14.5
## 5 NOEGMQ         37.1           -12.7
## 6 N103US         8.31           -12.9
```

Using measures of spread: sd(x), IQR(x), mad(x)

```
head(not_cancelled %>%
  group_by(dest) %>%
  summarize(distance_sd=sd(distance)) %>%
  arrange(desc(distance_sd)))
```

```
## # A tibble: 6 x 2
##   dest distance_sd
##   <chr>         <dbl>
## 1 EGE         10.5
## 2 SAN         10.3
## 3 SFO         10.2
## 4 HNL         10.0
## 5 SEA         9.98
## 6 LAS         9.91
```

Measures of rank: min(x), quantile(x, 0.25), max(x)

Question: When do the first and last flights leave each day?

```
head(not_cancelled %>%
  group_by(year, month, day) %>%
  summarise(
    first = min(dep_time),
    last = max (dep_time)
  ))
```

```
## # A tibble: 6 x 5
## # Groups:   year, month [1]
##   year month   day first  last
##   <int> <int> <int> <int> <int>
## 1  2013     1     1   517  2356
## 2  2013     1     2    42  2354
## 3  2013     1     3    32  2349
## 4  2013     1     4    25  2358
## 5  2013     1     5    14  2357
## 6  2013     1     6    16  2355
```

```
# Counts: n() -> returns the size of the current group
# sum(!is.na(x)) -> count the number of non-missing values
# n_distinct(x) -> count the number of distinct (unique) values
```

```
head(not_cancelled %>%
  group_by(dest) %>%
  summarise(carriers = n_distinct(carrier)) %>%
  arrange(desc(carriers)))
```

```
## # A tibble: 6 x 2
##   dest carriers
##   <chr>   <int>
## 1 ATL         7
## 2 BOS         7
## 3 CLT         7
## 4 ORD         7
## 5 TPA         7
## 6 AUS         6
```

```
head(not_cancelled %>%
  count(dest))
```

```
## # A tibble: 6 x 2
##   dest      n
##   <chr> <int>
## 1 ABQ    254
## 2 ACK    265
## 3 ALB   419
## 4 ANC     8
## 5 ATL 16898
## 6 AUS  2418
```

The total number of miles a plane flew:

```
head(not_cancelled %>%
  count(tailnum, wt =distance))
```

```
## # A tibble: 6 x 2
##   tailnum      n
##   <chr>    <dbl>
## 1 D942DN    3418
## 2 NOEGMQ   240626
## 3 N10156   110389
## 4 N102UW    25722
## 5 N103US    24619
## 6 N104UW    25157
```

How many flights left before 5am?

```
head(not_cancelled %>%
  group_by(year, month) %>%
  summarise(n_early = sum(dep_time < 500)))
```

```
## # A tibble: 6 x 3
## # Groups:   year [1]
##   year month n_early
##   <int> <int>   <int>
## 1  2013     1     75
## 2  2013     2     85
## 3  2013     3    147
## 4  2013     4    148
## 5  2013     5    120
## 6  2013     6    221
```

What proportion of flights are delayed by more than an hour in what month, the frequency, and which airlines?

```
not_cancelled %>%
  filter(!is.na(arr_delay)) %>%
  group_by(month) %>%
  summarise(hour_prop = mean(arr_delay > 60)) %>%
  arrange(desc(hour_prop))
```

```
## # A tibble: 12 x 2
##   month hour_prop
##   <int>    <dbl>
## 1     7    0.145
## 2     6    0.142
## 3    12    0.107
## 4     4    0.102
## 5     3    0.0837
## 6     8    0.0835
## 7     5    0.0795
## 8     1    0.0705
## 9     2    0.0689
## 10    9    0.0476
```

```
## 11    10    0.0451
## 12    11    0.0415

not_cancelled %>%
  filter(!is.na(arr_delay)) %>%
  group_by(carrier) %>%
  summarise(hour_prop = mean(arr_delay > 60), hour_freq = sum(arr_delay > 60)) %>%
  arrange(desc(hour_prop))
```

```
## # A tibble: 16 x 3
##   carrier hour_prop hour_freq
##   <chr>      <dbl>     <int>
## 1 00         0.138         4
## 2 YV         0.136        74
## 3 EV         0.133       6803
## 4 F9         0.128         87
## 5 FL         0.113        360
## 6 9E         0.106       1830
## 7 MQ         0.0928      2323
## 8 B6         0.0919      4965
## 9 WN         0.0883      1063
## 10 VX        0.0731        374
## 11 UA        0.0680      3931
## 12 AA        0.0648      2070
## 13 DL        0.0614      2927
## 14 US        0.0472         937
## 15 AS        0.0465         33
## 16 HA        0.0234          8
```

```
head(not_cancelled %>%
  filter(!is.na(arr_delay)) %>%
  group_by(carrier) %>%
  summarise(delay_freq = sum(arr_delay > 60)) %>%
  arrange(desc(delay_freq)))
```

```
## # A tibble: 6 x 2
##   carrier delay_freq
##   <chr>      <int>
## 1 EV         6803
## 2 B6         4965
## 3 UA         3931
## 4 DL         2927
## 5 MQ         2323
## 6 AA         2070
```

Rank airlines by the number of destinations that they fly to, considering only those airports that are flown to by two or more airlines.

```
flights %>%  
  # find all airports with > 1 carrier  
  group_by(dest) %>%  
  mutate(n_carriers = n_distinct(carrier)) %>%  
  # n_distinct is a faster and more concise equivalent of length(unique(x))  
  #- counts the number of unique values  
  
  filter(n_carriers > 1) %>%  
  
  # rank carriers by number of destinations  
  
  group_by(carrier) %>%  
  summarize(n_dest = n_distinct(dest)) %>%  
  arrange(desc(n_dest))
```

```
## # A tibble: 16 x 2  
##   carrier n_dest  
##   <chr>   <int>  
## 1 EV      51  
## 2 9E      48  
## 3 UA      42  
## 4 DL      39  
## 5 B6      35  
## 6 AA      19  
## 7 MQ      19  
## 8 WN      10  
## 9 OO       5  
## 10 US       5  
## 11 VX       4  
## 12 YV       3  
## 13 FL       2  
## 14 AS       1  
## 15 F9       1  
## 16 HA       1
```

*# What airline does the "EV" carrier code correspond to?*

```
filter(airlines, carrier == "EV" )
```

```
## # A tibble: 1 x 2  
##   carrier name  
##   <chr>   <chr>  
## 1 EV      ExpressJet Airlines Inc.
```

Find the 10 most delayed flights using a ranking function.


```
flights_delayed <- top_n(flights, 10, dep_delay)
select(flights_delayed, month, day, carrier, flight, dep_delay)
```

```
## # A tibble: 10 x 5
##   month   day carrier flight dep_delay
##   <int> <int> <chr>    <int>    <dbl>
## 1     1     9    HA         51    1301
## 2     1    10    MQ        3695    1126
## 3    12     5    AA         172     896
## 4     3    17    DL        2119     911
## 5     4    10    DL        2391     960
## 6     6    15    MQ        3535    1137
## 7     6    27    DL        2007     899
## 8     7    22    MQ        3075    1005
## 9     7    22    DL        2047     898
## 10    9    20    AA         177    1014
```



Show graphically the relationship between the distance and average delay for each location.

```
delays <- flights %>%  
  group_by(dest) %>%  
  summarise(  
    count=n(),  
    dist= mean(distance, na.rm=TRUE), delay = mean(arr_delay, na.rm=TRUE)  
  ) %>%  
  filter(count >20, dest != "HNL")  
  
ggplot(data= delays, mapping=aes(x=dist, y = delay)) +  
  geom_point(aes(size=count), alpha=1/3) +  
  geom_smooth(se=FALSE)
```




Delays are typically temporally correlated: even once the problem that caused the initial delay has been resolved, later flights are delayed to allow earlier flights to leave. Using `lag()` explore how the delay of a flight is related to the delay of the immediately preceding flight.

```
# This calculates the departure delay of the preceding flight from the same airport.

lagged_delays <- flights %>%
  arrange(origin, month, day, dep_time) %>%
  group_by(origin) %>%
  mutate(dep_delay_lag = lag(dep_delay)) %>%
  filter(!is.na(dep_delay), !is.na(dep_delay_lag))

# plots the relationship between the mean delay of a flight for all values of the previous flight.
# There seems to be an inverse "U" relationship
# between mean delay of a flight and the mean delay of the preceding flight.

lagged_delays %>%
  group_by(dep_delay_lag) %>%
  summarise(dep_delay_mean = mean(dep_delay)) %>%
  ggplot(aes(y= dep_delay_mean, x=dep_delay_lag)) +
  geom_point() +
  scale_x_continuous(breaks = seq(0, 1500, by = 120)) +
  labs(y = "Departure Delay", x = "Previous Departure Delay")
```



```
# The overall relationship looks similar in all three origin airports.
```

```
lagged_delays %>%  
  group_by(origin, dep_delay_lag) %>%  
  summarise(dep_delay_mean = mean(dep_delay)) %>%  
  ggplot(aes(y = dep_delay_mean, x = dep_delay_lag)) +  
  geom_point() +  
  facet_wrap(~ origin, ncol=1) +  
  labs(y = "Departure Delay", x = "Previous Departure Delay")
```

