

# Text Mining

Monica Buczynski

October 09, 2020

Note: The purpose of this document is to showcase a sample of skills that I learned in *Text Mining with R: A Tidy Approach* by Julia Silge and David Robinson. Some scripts were taken from <https://www.tidytextmining.com/s.html>. The code for each exercise was studied carefully for understanding and then was retyped manually into R to maximize the learning experience; however, many of the scripts were altered for further analysis and presentation aesthetics. Additionally, I added my own code for further analysis and my own curiosity.

Skills that I focused on included:

- The tidy text format
- Sentiment analysis with tidy data
- Analyzing word and document frequency: tf-idf
- Relationships between words: n-grams and correlations
- Converting to and from non-tidy formats

# 1 The tidy text format

## 1.2 The unnest\_tokens function

```
text <- c("Because I could not stop for Death -",
         "He kindly stopped for me -",
         "The Carriage held but just Ourselves -",
         "and Immortality")
```

```
text
```

```
## [1] "Because I could not stop for Death -"
## [2] "He kindly stopped for me -"
## [3] "The Carriage held but just Ourselves -"
## [4] "and Immortality"
```

```
# save as df
```

```
library(dplyr)
text_df <- tibble(line = 1:4, text = text)
text_df
```

```
## # A tibble: 4 x 2
##   line text
##   <int> <chr>
## 1     1 Because I could not stop for Death -
## 2     2 He kindly stopped for me -
## 3     3 The Carriage held but just Ourselves -
## 4     4 and Immortality
```

```
# tokenization
```

```
library(tidytext)

text_df %>%
  unnest_tokens(word, text, to_lower = FALSE)
```

```
## # A tibble: 20 x 2
##   line word
##   <int> <chr>
## 1     1 Because
## 2     1 I
## 3     1 could
## 4     1 not
## 5     1 stop
## 6     1 for
## 7     1 Death
## 8     2 He
## 9     2 kindly
## 10    2 stopped
## 11    2 for
## 12    2 me
## 13    3 The
## 14    3 Carriage
## 15    3 held
## 16    3 but
```

```
## 17      3 just
## 18      3 Ourselves
## 19      4 and
## 20      4 Immortality
```

```
# do not use *to_lower = FALSE* to convert the tokens to lowercase
```

```
text_df %>%
  unnest_tokens(word, text)
```

```
## # A tibble: 20 x 2
##   line word
##   <int> <chr>
## 1     1 because
## 2     1 i
## 3     1 could
## 4     1 not
## 5     1 stop
## 6     1 for
## 7     1 death
## 8     2 he
## 9     2 kindly
## 10    2 stopped
## 11    2 for
## 12    2 me
## 13    3 the
## 14    3 carriage
## 15    3 held
## 16    3 but
## 17    3 just
## 18    3 ourselves
## 19    4 and
## 20    4 immortality
```

### 1.3 Tidying the works of Jane Austen

```
original_books <- austen_books() %>%
  group_by(book) %>%
  mutate(linenumber = row_number(),
         chapter = cumsum(str_detect(text, regex("^chapter [\\divxlc]",
                                                ignore_case = TRUE)))) %>% ungroup()
```

original\_books

```
## # A tibble: 73,422 x 4
##   text                book          linenumber chapter
##   <chr>              <fct>          <int>    <int>
## 1 "SENSE AND SENSIBILITY" Sense & Sensibility      1      0
## 2 ""                Sense & Sensibility      2      0
## 3 "by Jane Austen"   Sense & Sensibility      3      0
## 4 ""                Sense & Sensibility      4      0
## 5 "(1811)"          Sense & Sensibility      5      0
## 6 ""                Sense & Sensibility      6      0
## 7 ""                Sense & Sensibility      7      0
## 8 ""                Sense & Sensibility      8      0
## 9 ""                Sense & Sensibility      9      0
## 10 "CHAPTER 1"       Sense & Sensibility     10      1
## # ... with 73,412 more rows
```

```
# restructure df in the one-token-per-row format with the unnest_tokens()
```

```
tidy_books <- original_books %>%
  unnest_tokens(words, text)
```

tidy\_books

```
## # A tibble: 725,055 x 4
##   book          linenumber chapter words
##   <fct>          <int>    <int> <chr>
## 1 Sense & Sensibility      1      0 sense
## 2 Sense & Sensibility      1      0 and
## 3 Sense & Sensibility      1      0 sensibility
## 4 Sense & Sensibility      3      0 by
## 5 Sense & Sensibility      3      0 jane
## 6 Sense & Sensibility      3      0 austen
## 7 Sense & Sensibility      5      0 1811
## 8 Sense & Sensibility     10      1 chapter
## 9 Sense & Sensibility     10      1 1
## 10 Sense & Sensibility     13      1 the
## # ... with 725,045 more rows
```

```
# add stop words - words that are not usefull to us for analysis
```

stop\_words

```
## # A tibble: 1,149 x 2
##   word      lexicon
##   <chr>    <chr>
## 1 a       SMART
```

```
## 2 a's          SMART
## 3 able         SMART
## 4 about        SMART
## 5 above        SMART
## 6 according    SMART
## 7 accordingly  SMART
## 8 across       SMART
## 9 actually     SMART
## 10 after       SMART
## # ... with 1,139 more rows

# Practice adding a new row
newRow <- data.frame(word="AAAAA",lexicon = "SMART" )
stop_words <- rbind(stop_words, newRow)

tidy_books <- tidy_books %>%
  rename("word" = "words") %>% # rename column name "words" to "word" in tidy_books
                                # so that there is a key between tidy_books and
                                # stop_words for anti_join()
  anti_join(stop_words, by = "word") # drops all observations in x that have a match in y

tidy_books

## # A tibble: 217,609 x 4
##   book          linenumber chapter word
##   <fct>          <int>     <int> <chr>
## 1 Sense & Sensibility      1         0 sense
## 2 Sense & Sensibility      1         0 sensibility
## 3 Sense & Sensibility      3         0 jane
## 4 Sense & Sensibility      3         0 austen
## 5 Sense & Sensibility      5         0 1811
## 6 Sense & Sensibility     10         1 chapter
## 7 Sense & Sensibility     10         1 1
## 8 Sense & Sensibility     13         1 family
## 9 Sense & Sensibility     13         1 dashwood
## 10 Sense & Sensibility     13         1 settled
## # ... with 217,599 more rows

# use count() to find the most common words

tidy_books %>%
  count(word, sort = TRUE)

## # A tibble: 13,914 x 2
##   word      n
##   <chr> <int>
## 1 miss   1855
## 2 time   1337
## 3 fanny   862
## 4 dear    822
## 5 lady    817
## 6 sir     806
## 7 day     797
## 8 emma    787
## 9 sister  727
```

```
## 10 house      699
## # ... with 13,904 more rows
```

## 1.4 The gutenbergr package

```
BooksOz <- gutenbergr_metadata[grep("Oz", gutenbergr_metadata$title), ]
BooksOz
```

```
## # A tibble: 49 x 8
##   gutenbergr_id title   author   gutenbergr_autho~ language gutenbergr_books~ rights
##         <int> <chr>   <chr>         <int> <chr>         <chr>         <chr>
## 1           54 The M~ Baum, ~         42 en      Children's Lite~ Publi~
## 2           55 The W~ Baum, ~         42 en      Children's Lite~ Publi~
## 3          419 The M~ Baum, ~         42 en      Children's Lite~ Publi~
## 4          420 Dorot~ Baum, ~         42 en      Children's Lite~ Publi~
## 5          485 The R~ Baum, ~         42 en      Children's Lite~ Publi~
## 6          486 Ozma ~ Baum, ~         42 en      Fantasy/Childre~ Publi~
## 7          517 The E~ Baum, ~         42 en      Children's Lite~ Publi~
## 8          955 The P~ Baum, ~         42 en      Children's Lite~ Publi~
## 9          956 Tik-T~ Baum, ~         42 en      Children's Lite~ Publi~
## 10         957 The S~ Baum, ~         42 en      Children's Lite~ Publi~
## # ... with 39 more rows, and 1 more variable: has_text <lgl>
```

```
#gutenbergr_metadata %>%
#filter(title == "Oz")
```

```
WWOz <- gutenbergr_download(55)
```

```
WWOz
```

```
## # A tibble: 4,750 x 2
##   gutenbergr_id text
##         <int> <chr>
## 1           55 "[Illustration]"
## 2           55 ""
## 3           55 ""
## 4           55 ""
## 5           55 ""
## 6           55 "The Wonderful Wizard of Oz"
## 7           55 ""
## 8           55 "by L. Frank Baum"
## 9           55 ""
## 10          55 ""
## # ... with 4,740 more rows
```

```
tidy_books_Oz <- WWOz %>%
  unnest_tokens(words, text)
```

```
tidy_books_Oz
```

```
## # A tibble: 39,765 x 2
##   gutenbergr_id words
##         <int> <chr>
## 1           55 illustration
## 2           55 the
## 3           55 wonderful
## 4           55 wizard
## 5           55 of
## 6           55 oz
```

```
## 7          55 by
## 8          55 l
## 9          55 frank
## 10         55 baum
## # ... with 39,755 more rows

tidy_books_Oz <- tidy_books_Oz %>%
  rename("word" = "words") %>%
  # rename column name "words" to "word" in tidy_books so that there is a key
  # between tidy_books and stop_words for anti_join()
anti_join(stop_words, by = "word")
# drops all observations in x that have a match in y

tidy_books_Oz %>%
  count(word, sort = TRUE) %>%
  summary(tidy_books_Oz$n)
```

```
##      word          n
## Length:2533      Min.   : 1.000
## Class :character 1st Qu.: 1.000
## Mode  :character Median : 2.000
##                  Mean   : 4.941
##                  3rd Qu.: 4.000
##                  Max.   :347.000
```

```
tidy_books_Oz %>%
  count(word, sort = TRUE)
```

```
## # A tibble: 2,533 x 2
##   word          n
##   <chr>      <int>
## 1 dorothy    347
## 2 scarecrow 219
## 3 woodman    176
## 4 lion       173
## 5 oz         164
## 6 tin        140
## 7 witch      125
## 8 green      104
## 9 girl        93
## 10 head       90
## # ... with 2,523 more rows
```

```
tidy_books_Oz %>%
  count(word, sort = TRUE) %>%
  dplyr::filter(word == "munchkins")
```

```
## # A tibble: 1 x 2
##   word          n
##   <chr>      <int>
## 1 munchkins    21
```

```
tidy_books_Oz %>%
  count(word, sort = TRUE) %>%
  dplyr::filter(word == "monkeys")
```

```
## # A tibble: 1 x 2
```



```
## word n
## <chr> <int>
## 1 monkeys 44
```

```
library("RColorBrewer")
display.brewer.pal(n = 8, name = 'Dark2')
```



Dark2 (qualitative)

Figure 1: Words mentioned more than 50 times in The Wonderful Wizard of Oz

```
brewer.pal(n = 8, name = 'Dark2')
```

```
## [1] "#1B9E77" "#D95F02" "#7570B3" "#E7298A" "#66A61E" "#E6AB02" "#A6761D"
## [8] "#666666"
```

```
tidy_books_Oz %>%
  count(word, sort= TRUE) %>%
  dplyr::filter(n >50) %>%
  mutate(word = reorder(word, n)) %>%
  ggplot(aes(word, n)) +
  geom_col(fill= "#1B9E77"
) +
  labs(y = "frequency") +
  coord_flip()
```




Figure 2: Words mentioned more than 50 times in The Wonderful Wizard of Oz

## 2 Sentiment analysis with tidy data

### 2.1 The sentiments dataset

```
get_sentiments("afinn")
```

```
## # A tibble: 2,477 x 2
##   word      value
##   <chr>    <dbl>
## 1 abandon     -2
## 2 abandoned   -2
## 3 abandons    -2
## 4 abducted    -2
## 5 abduction   -2
## 6 abductions  -2
## 7 abhor       -3
## 8 abhorred    -3
## 9 abhorrent   -3
## 10 abhors     -3
## # ... with 2,467 more rows
```

```
get_sentiments("nrc")
```

```
## # A tibble: 13,901 x 2
##   word      sentiment
##   <chr>    <chr>
## 1 abacus      trust
## 2 abandon     fear
## 3 abandon     negative
## 4 abandon     sadness
## 5 abandoned   anger
## 6 abandoned   fear
## 7 abandoned   negative
## 8 abandoned   sadness
## 9 abandonment anger
## 10 abandonment fear
## # ... with 13,891 more rows
```

```
get_sentiments("bing")
```

```
## # A tibble: 6,786 x 2
##   word      sentiment
##   <chr>    <chr>
## 1 2-faces   negative
## 2 abnormal  negative
## 3 abolish   negative
## 4 abominable negative
## 5 abominably negative
## 6 abominate  negative
## 7 abomination negative
## 8 abort      negative
## 9 aborted    negative
## 10 abortions negative
## # ... with 6,776 more rows
```

## 2.2 Sentiment analysis with inner join

```
tidy_books <- austen_books() %>%
  group_by(book) %>%
  mutate(linenumber = row_number(),
         chapter = cumsum(str_detect(text, regex("^chapter [\\divxlc]",
                                                ignore_case = TRUE)))) %>%
  ungroup() %>%
  unnest_tokens(word, text)

# What are the most common joy words in Emma?

nrc_joy <- get_sentiments("nrc") %>%
  filter(sentiment == "joy")

tidy_books %>%
  filter(book == "Emma") %>%
  inner_join(nrc_joy, by = "word") %>%
  count(word, sort = TRUE)

## # A tibble: 303 x 2
##   word      n
##   <chr>  <int>
## 1 good    359
## 2 young   192
## 3 friend  166
## 4 hope    143
## 5 happy   125
## 6 love    117
## 7 deal     92
## 8 found    92
## 9 present  89
## 10 kind    82
## # ... with 293 more rows

# examine how sentiment changes throughout each novel

jane_austen_sentiment <- tidy_books %>%
  inner_join(get_sentiments("bing"), by = "word") %>%
  count(book, index = linenumber %/% 80, sentiment) %>%
  pivot_wider(names_from = sentiment, values_from = n, values_fill = 0) %>%
  mutate(sentiment = positive - negative)
```

```
ggplot(jane_austen_sentiment, aes(index, sentiment, fill = book)) +
  geom_col(show.legend = FALSE) +
  facet_wrap(~book, ncol = 2, scales = "free_x") +
  scale_fill_brewer(palette = "Dark2")
```

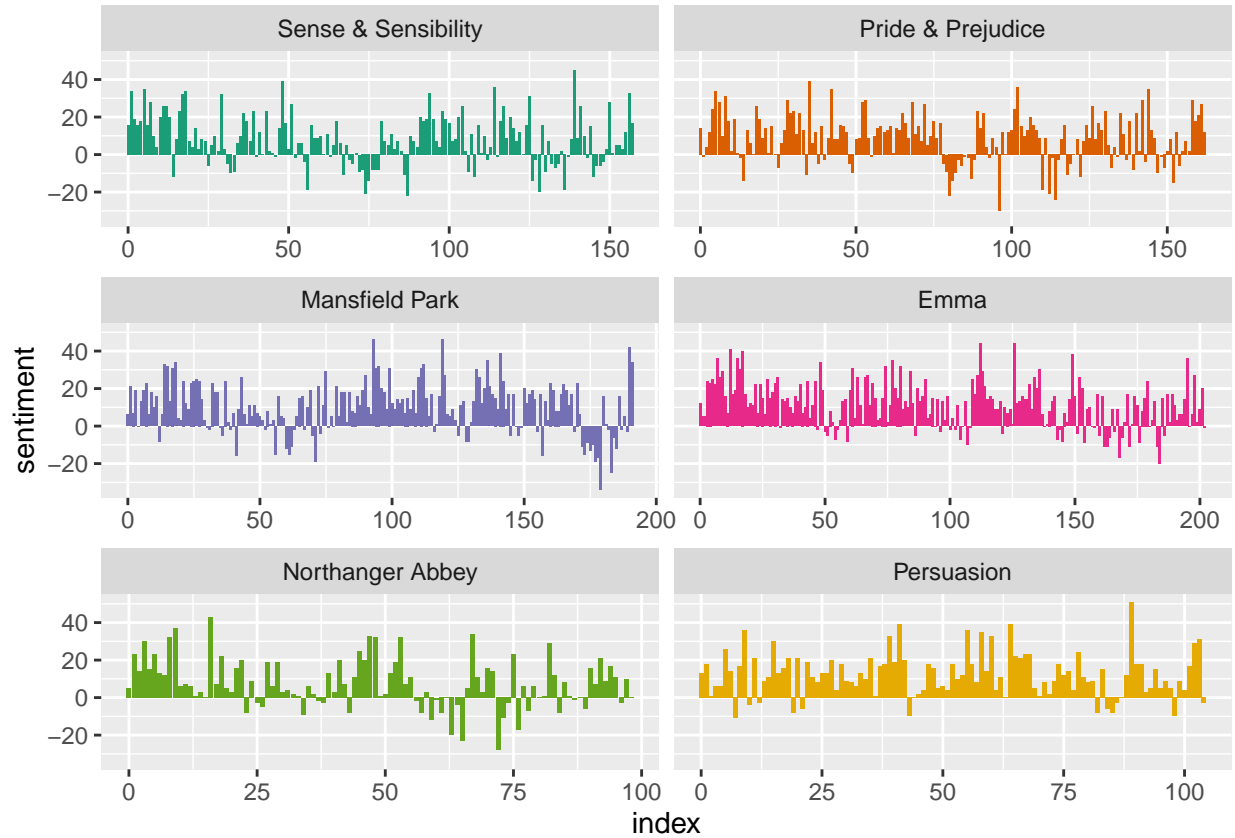


Figure 3: Sentiment through the narratives of Jane Austen's novels

## 2.3 Comparing the three sentiment dictionaries

```
# filtering to one novel that I am interested in

pride_prejudice <- tidy_books %>%
  filter(book == "Pride & Prejudice")

pride_prejudice

## # A tibble: 122,204 x 4
##   book          linenumber chapter word
##   <fct>          <int>    <int> <chr>
## 1 Pride & Prejudice      1      0 pride
## 2 Pride & Prejudice      1      0 and
## 3 Pride & Prejudice      1      0 prejudice
## 4 Pride & Prejudice      3      0 by
## 5 Pride & Prejudice      3      0 jane
## 6 Pride & Prejudice      3      0 austen
## 7 Pride & Prejudice      7      1 chapter
## 8 Pride & Prejudice      7      1 1
## 9 Pride & Prejudice     10      1 it
## 10 Pride & Prejudice     10      1 is
## # ... with 122,194 more rows

# need two different patterns because AFINN has a numeric measure while Bing and NRC are binary.

afinn <- pride_prejudice %>%
  inner_join(get_sentiments("afinn"), by = "word") %>%
  group_by(index = linenumber %/% 80) %>%
  summarise(sentiment = sum(value)) %>%
  mutate(method = "AFINN")

bing_and_nrc <- bind_rows(pride_prejudice %>%
  inner_join(get_sentiments("bing"), by = "word") %>%
  mutate(method = "Bind et al."),
  pride_prejudice %>%
  inner_join(get_sentiments("nrc"), by = "word") %>%
  filter(sentiment %in% c("positive", "negative")) %>%
  mutate(method = "NRC")) %>%
  count(method, index = linenumber %/% 80, sentiment) %>%
  pivot_wider(names_from = sentiment, values_from = n, values_fill = 0) %>%
  mutate(sentiment = positive - negative)
```

```
bind_rows(afinn, bing_and_nrc) %>%
  ggplot(aes(x = index, y = sentiment, fill = method)) +
  geom_col(show.legend = FALSE) +
  facet_wrap(~method, ncol = 1, scales = "free_y") +
  scale_fill_brewer(palette = "Dark2")
```




Figure 4: Comparing three sentiment lexicons using Pride and Prejudice



Why is, for example, the result for the NRC lexicon biased so high in sentiment compared to the Bing et al. result? Both lexicons have more negative than positive words, but the ratio of negative to positive words is higher in the Bing lexicon than the NRC lexicon.

```
nrc <- get_sentiments("nrc") %>%
  filter(sentiment %in% c("positive", "negative")) %>%
  count(sentiment) %>%
  mutate(proportion = n/sum(n),
         lexicon = "NRC")

bing <- get_sentiments("bing") %>%
  count(sentiment) %>%
  mutate(proportion = n/sum(n),
         lexicon = "bing")

full_join(bing, nrc)
```

```
## # A tibble: 4 x 4
##   sentiment      n proportion lexicon
##   <chr>      <int>      <dbl> <chr>
## 1 negative   4781      0.705 bing
## 2 positive   2005      0.295 bing
## 3 negative   3324      0.590 NRC
## 4 positive   2312      0.410 NRC
```

## 2.5 Wordclouds

Illustrate the most common words in Jane Austen's works in a word cloud

```
tidy_books %>%
  anti_join(stop_words) %>%
  count(word) %>%
  with(wordcloud(word, n, max.words = 50))
```



```
library(reshape2)

tidy_books %>%
  inner_join(get_sentiments("bing")) %>%
  count(word, sentiment, sort = TRUE) %>%
  acast(word ~ sentiment, value.var = "n", fill = 0) %>%
  comparison.cloud(colors = c("gray20", "gray80"),
    max.words = 100)
```

negative



positive

## 2.6 Looking at units beyond just words

```
# tokenizing at the sentence level

PandP_sentences <- tibble(text = prideprejudice) %>%
  unnest_tokens(sentence, text, token = "sentences")

# look at sentence #2

PandP_sentences$sentence[2]

## [1] "by jane austen"

# tokenizing at the chapter level

austen_chapters <- austen_books() %>%
  group_by(book) %>%
  unnest_tokens(chapter, text, token = "regex",
    pattern = "Chapter|CHAPTER [\\dIVXLC]") %>% ungroup()

austen_chapters %>%
  group_by(book) %>%
  summarise(chapters = n())

## # A tibble: 6 x 2
##   book          chapters
##   <fct>          <int>
## 1 Sense & Sensibility      51
## 2 Pride & Prejudice       62
## 3 Mansfield Park         49
## 4 Emma                  56
## 5 Northanger Abbey       32
## 6 Persuasion             25

## What are the most negative chapters in each of Jane Austen's novels?

bingnegative <- get_sentiments("bing") %>%
  filter(sentiment == "negative")

wordcounts <- tidy_books %>%
  group_by(book, chapter) %>%
  summarize(words = n())

tidy_books %>%
  semi_join(bingnegative) %>%
  group_by(book, chapter) %>%
  summarize(negativewords = n()) %>%
  left_join(wordcounts, by = c("book", "chapter")) %>%
  mutate(ratio = negativewords/words) %>%
  filter(chapter != 0) %>%
  top_n(1) %>%
  ungroup()

## # A tibble: 6 x 5
##   book          chapter negativewords words  ratio
##   <fct>          <int>          <int> <int>  <dbl>
```

|                          |    |     |      |        |
|--------------------------|----|-----|------|--------|
| ## 1 Sense & Sensibility | 43 | 161 | 3405 | 0.0473 |
| ## 2 Pride & Prejudice   | 34 | 111 | 2104 | 0.0528 |
| ## 3 Mansfield Park      | 46 | 173 | 3685 | 0.0469 |
| ## 4 Emma                | 15 | 151 | 3340 | 0.0452 |
| ## 5 Northanger Abbey    | 21 | 149 | 2982 | 0.0500 |
| ## 6 Persuasion          | 4  | 62  | 1807 | 0.0343 |

### 3 Analyzing word and document frequency: tf-idf

The statistic tf-idf is intended to measure how important a word is to a document in a collection (or corpus) of documents, for example, to one novel in a collection of novels or to one website in a collection of websites.

For a term  $t$  in a document  $d$ , the weight  $W_{t,d}$  of term  $t$  in document  $d$  is given by:

$$tf-idf W_{t,d} = TF_t \cdot \log(N/DF_t)$$

Where:

$TF_{t,d}$  is the number of occurrences of  $t$  in document  $d$ .  $DF_t$  is the number of documents containing the term  $t$ .  $N$  is the total number of documents in the corpus.

TF  $\rightarrow$  term frequency IDF  $\rightarrow$  inverse document frequency - decreases the weight for commonly used words and increases the weight for words that are not used very much in a collection of documents

The higher the TF\*IDF score (weight), the rarer the term and vice versa

### 3.1 Term frequency in Jane Austen's novels

What are the most commonly used words in Jane Austen's novels?

```
book_words <- austen_books() %>%
  unnest_tokens(word, text) %>%
  count(book, word, sort = TRUE)

total_words <- book_words %>%
  group_by(book) %>%
  summarize(total = sum(n))

book_words <- full_join(book_words, total_words)

book_words
```

```
## # A tibble: 40,379 x 4
##   book          word      n total
##   <fct>         <chr> <int> <int>
## 1 Mansfield Park the      6206 160460
## 2 Mansfield Park to       5475 160460
## 3 Mansfield Park and       5438 160460
## 4 Emma          to       5239 160996
## 5 Emma          the       5201 160996
## 6 Emma          and       4896 160996
## 7 Mansfield Park of       4778 160460
## 8 Pride & Prejudice the     4331 122204
## 9 Emma          of       4291 160996
## 10 Pride & Prejudice to     4162 122204
## # ... with 40,369 more rows
```

```
ggplot(book_words, aes(n/total, fill = book)) +
  geom_histogram(show.legend = FALSE) +
  xlim(NA, 0.0009) +
  facet_wrap(~book, ncol = 2, scales = "free_y") +
  scale_fill_brewer(palette = "Dark2")
```

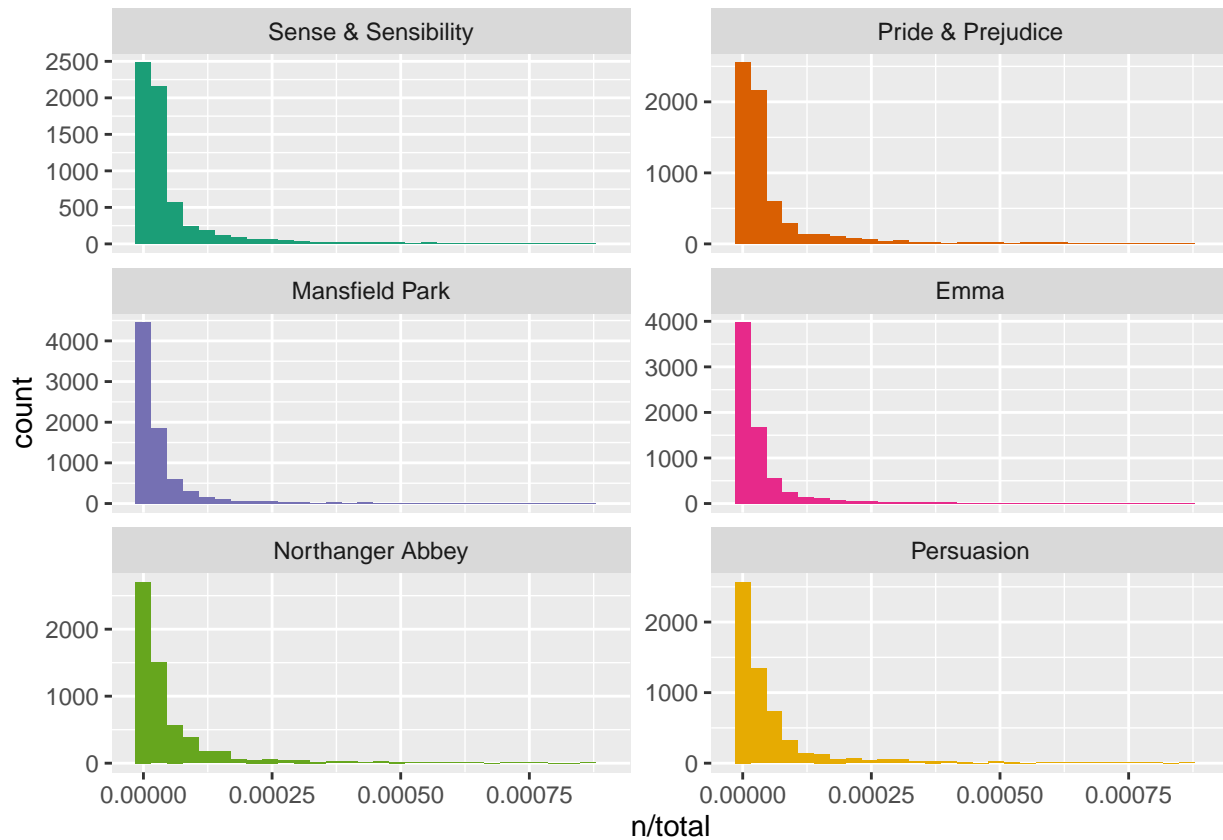


Figure 5: Term Frequency Distribution in Jane Austen's Novels

*# Observation: many words that occur rarely and fewer words that occur frequently*



## 3.2 Zipf's law

Zipf's law states that the frequency that a word appears is inversely proportional to its rank.

```
freq_by_rank <- book_words %>%  
  group_by(book) %>%  
  mutate(rank = row_number(),  
         "term_frequency" = n/total)
```

```
freq_by_rank
```

```
## # A tibble: 40,379 x 6  
## # Groups:   book [6]  
##   book          word      n  total  rank term_frequency  
##   <fct>         <chr> <int> <int> <int>         <dbl>  
## 1 Mansfield Park the      6206 160460     1         0.0387  
## 2 Mansfield Park to       5475 160460     2         0.0341  
## 3 Mansfield Park and      5438 160460     3         0.0339  
## 4 Emma          to       5239 160996     1         0.0325  
## 5 Emma          the      5201 160996     2         0.0323  
## 6 Emma          and      4896 160996     3         0.0304  
## 7 Mansfield Park of       4778 160460     4         0.0298  
## 8 Pride & Prejudice the    4331 122204     1         0.0354  
## 9 Emma          of       4291 160996     4         0.0267  
## 10 Pride & Prejudice to    4162 122204     2         0.0341  
## # ... with 40,369 more rows
```

```
freq_by_rank %>%
  ggplot(aes(x = rank, y = term_frequency, color = book)) + geom_line(size = 1.1, alpha = 0.8, show.legend = FALSE) +
  scale_y_log10() +
  labs(y = "term frequency")
```




Figure 6: Zipf's law for Jane Austen's novels

```
# rank column tells the rank of each word within the frequency table
```

### 3.3 The bind\_tf\_idf function

```
book_words <- book_words %>%  
  bind_tf_idf(word, book, n)
```

```
book_words
```

```
## # A tibble: 40,379 x 7  
##   book          word      n total    tf    idf tf_idf  
##   <fct>        <chr> <int> <int> <dbl> <dbl> <dbl>  
## 1 Mansfield Park the      6206 160460 0.0387 0 0  
## 2 Mansfield Park to      5475 160460 0.0341 0 0  
## 3 Mansfield Park and     5438 160460 0.0339 0 0  
## 4 Emma         to      5239 160996 0.0325 0 0  
## 5 Emma         the     5201 160996 0.0323 0 0  
## 6 Emma         and     4896 160996 0.0304 0 0  
## 7 Mansfield Park of      4778 160460 0.0298 0 0  
## 8 Pride & Prejudice the    4331 122204 0.0354 0 0  
## 9 Emma         of      4291 160996 0.0267 0 0  
## 10 Pride & Prejudice to    4162 122204 0.0341 0 0  
## # ... with 40,369 more rows
```

```
book_words %>%  
  select(-total) %>%  
  arrange(desc(tf_idf))
```

```
## # A tibble: 40,379 x 6  
##   book          word      n    tf    idf tf_idf  
##   <fct>        <chr> <int> <dbl> <dbl> <dbl>  
## 1 Sense & Sensibility elinor     623 0.00519 1.79 0.00931  
## 2 Sense & Sensibility marianne  492 0.00410 1.79 0.00735  
## 3 Mansfield Park      crawford  493 0.00307 1.79 0.00551  
## 4 Pride & Prejudice    darcy    373 0.00305 1.79 0.00547  
## 5 Persuasion          elliot   254 0.00304 1.79 0.00544  
## 6 Emma               emma     786 0.00488 1.10 0.00536  
## 7 Northanger Abbey    tilney   196 0.00252 1.79 0.00452  
## 8 Emma               weston   389 0.00242 1.79 0.00433  
## 9 Pride & Prejudice    bennet   294 0.00241 1.79 0.00431  
## 10 Persuasion          wentworth 191 0.00228 1.79 0.00409  
## # ... with 40,369 more rows
```

```

book_words %>%
  arrange(desc(tf_idf)) %>%
  mutate(word = factor(word, levels = rev(unique(word)))) %>%
  group_by(book) %>%
  top_n(15) %>%
  ungroup() %>%
  ggplot(aes(x= word, y= tf_idf, fill = book)) +
  geom_col(show.legend = FALSE) +
  labs(x = NULL, y = "tf-idf") +
  facet_wrap(~book, ncol = 2, scales = "free") +
  coord_flip() +
  scale_fill_brewer(palette = "Dark2")

```

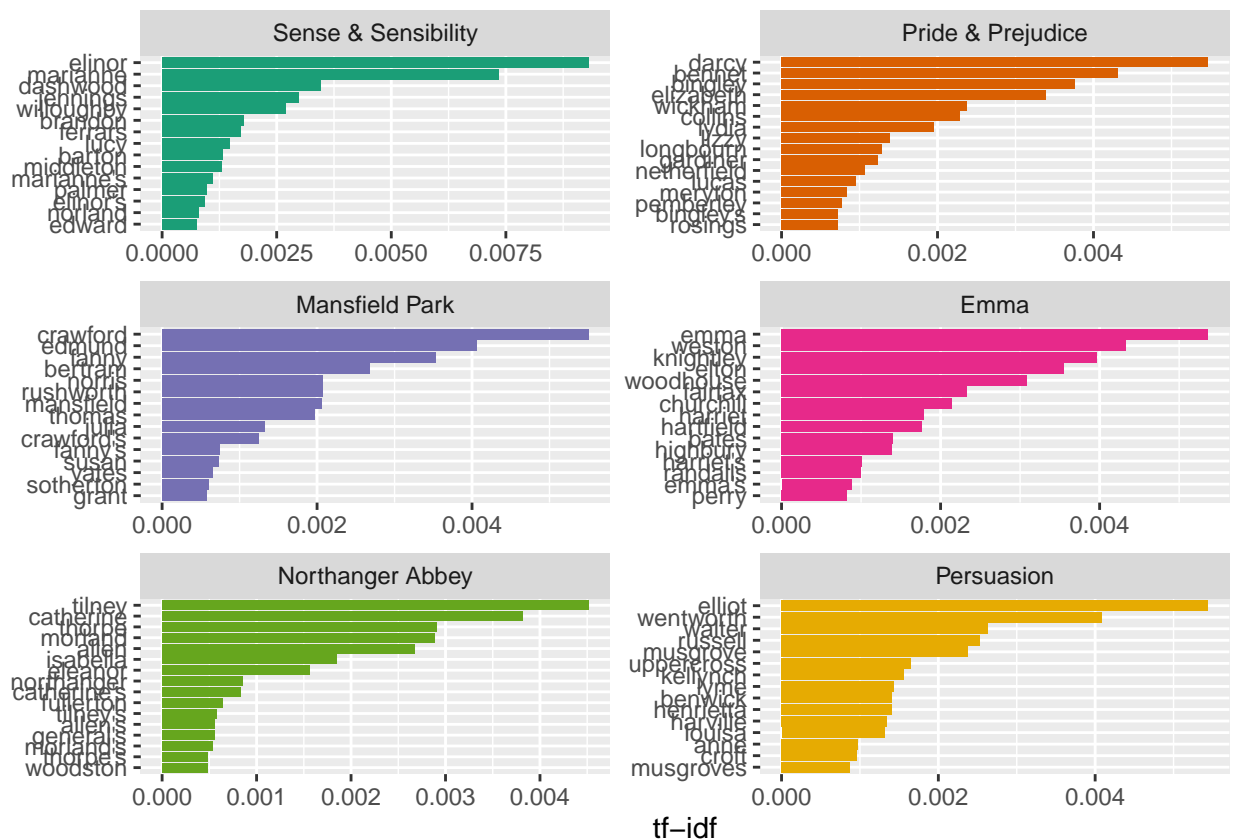


Figure 7: Highest tf-idf words in each of Jane Austen's Novels

## 4.2 Counting and correlating pairs of words with the widyr package

### 4.2.1 Counting and correlating among sections

The widyr package makes operations such as computing counts and correlations easy, by simplifying the pattern of “widen data, perform an operation, then re-tidy data”. The book “Pride and Prejudice” divided into 10-line sections, as we did (with larger sections) for sentiment analysis in Chapter 2. We may be interested in what words tend to appear within the same section.

```
austen_section_words <- austen_books() %>%
  filter(book == "Pride & Prejudice" ) %>%
  mutate(section = row_number() %/% 10) %>%
  filter(section > 0) %>%
  unnest_tokens(word, text) %>%
  filter(!word %in% stop_words$word)
```

```
austen_section_words
```

```
## # A tibble: 37,240 x 3
##   book          section word
##   <fct>          <dbl> <chr>
## 1 Pride & Prejudice      1 truth
## 2 Pride & Prejudice      1 universally
## 3 Pride & Prejudice      1 acknowledged
## 4 Pride & Prejudice      1 single
## 5 Pride & Prejudice      1 possession
## 6 Pride & Prejudice      1 fortune
## 7 Pride & Prejudice      1 wife
## 8 Pride & Prejudice      1 feelings
## 9 Pride & Prejudice      1 views
## 10 Pride & Prejudice     1 entering
## # ... with 37,230 more rows
```

```
# count words co-occurring within sections
```

```
word_pairs <- austen_section_words %>% pairwise_count(word, section, sort = TRUE)
```

```
word_pairs
```

```
## # A tibble: 796,008 x 3
##   item1    item2      n
##   <chr>    <chr>  <dbl>
## 1 darcy    elizabeth  144
## 2 elizabeth darcy    144
## 3 miss     elizabeth  110
## 4 elizabeth miss    110
## 5 elizabeth jane    106
## 6 jane     elizabeth  106
## 7 miss     darcy     92
## 8 darcy    miss      92
## 9 elizabeth bingley   91
## 10 bingley  elizabeth  91
## # ... with 795,998 more rows
```

```
# the most common pair of words in a section is "Elizabeth" and "Darcy"
# (the two main characters)
```

```
word_pairs %>%
  filter(item1 == "darcy")
```

```
## # A tibble: 2,930 x 3
##   item1 item2      n
##   <chr> <chr>   <dbl>
## 1 darcy elizabeth 144
## 2 darcy miss      92
## 3 darcy bingley   86
## 4 darcy jane      46
## 5 darcy bennet    45
## 6 darcy sister    45
## 7 darcy time      41
## 8 darcy lady      38
## 9 darcy friend    37
## 10 darcy wickham   37
## # ... with 2,920 more rows
```

#### 4.2.2 Pairwise correlations

Examine among words, which indicates how often they appear together relative to how often they appear separately.

```
#filter for common words first
word_cors <- austen_section_words %>%
  group_by(word) %>%
  filter(n() >= 20) %>%
  pairwise_cor(word, section) %>%
  arrange(correlation)
```

```
word_cors
```

```
## # A tibble: 154,842 x 3
##   item1      item2 correlation
##   <chr>   <chr>         <dbl>
## 1 darcy   lydia         -0.122
## 2 lydia   darcy         -0.122
## 3 collins bingley       -0.122
## 4 bingley collins       -0.122
## 5 jane    lady          -0.111
## 6 lady    jane          -0.111
## 7 collins darcy         -0.100
## 8 darcy    collins       -0.100
## 9 longbourn darcy        -0.0946
## 10 darcy    longbourn     -0.0946
## # ... with 154,832 more rows
```

```
word_cors <- austen_section_words %>%
  group_by(word) %>%
  filter(n() >= 20) %>%
  pairwise_cor(word, section) %>%
  arrange(desc(correlation))
```

```
word_cors
```

```
## # A tibble: 154,842 x 3
```

```
##   item1   item2   correlation
##   <chr>   <chr>       <dbl>
## 1 bourgh   de         0.951
## 2 de       bourgh     0.951
## 3 pounds   thousand   0.701
## 4 thousand pounds     0.701
## 5 william   sir        0.664
## 6 sir       william    0.664
## 7 catherine lady      0.663
## 8 lady      catherine  0.663
## 9 forster   colonel    0.622
## 10 colonel  forster     0.622
## # ... with 154,832 more rows
```


*# find the words most correlated with a word like "pounds" using a filter operation.*

```
word_cors %>%
  filter(item1 == "pounds") %>%
  arrange(desc(correlation))
```

```
## # A tibble: 393 x 3
##   item1 item2   correlation
##   <chr> <chr>       <dbl>
## 1 pounds thousand   0.701
## 2 pounds ten        0.231
## 3 pounds fortune    0.164
## 4 pounds settled    0.149
## 5 pounds wickham's  0.142
## 6 pounds children   0.129
## 7 pounds mother's   0.119
## 8 pounds believed   0.0932
## 9 pounds estate     0.0890
## 10 pounds ready     0.0860
## # ... with 383 more rows
```

*# pick particular interesting words and find the other words most associated with them*

```
word_cors %>%
  filter(item1 %in% c("elizabeth", "pounds", "married", "pride")) %>%
  group_by(item1) %>%
  top_n(6) %>%
  ungroup() %>%
  mutate(item2 = reorder(item2, correlation)) %>%
  ggplot(aes(x = item2, y = correlation, fill = item1)) +
  geom_bar(stat = "identity", show.legend = FALSE) +
  facet_wrap(~ item1, scales = "free") +
  labs(x = "item 2") +
  coord_flip() +
  scale_fill_brewer(palette = "Dark2")
```





## Converting to and from non-tidy formats

### Tidying dfm objects

```
data("data_corpus_inaugural", package = "quanteda")

inaug_dfm <- quanteda::dfm(data_corpus_inaugural, verbose = FALSE)

# dfm = document-feature-matrix

# integrate into tidy

inaug_td <- tidy(inaug_dfm)
inaug_td
```

```
## # A tibble: 45,453 x 3
##   document      term      count
##   <chr>         <chr>    <dbl>
## 1 1789-Washington fellow-citizens 1
## 2 1797-Adams     fellow-citizens 3
## 3 1801-Jefferson fellow-citizens 2
## 4 1809-Madison   fellow-citizens 1
## 5 1813-Madison   fellow-citizens 1
## 6 1817-Monroe    fellow-citizens 5
## 7 1821-Monroe    fellow-citizens 1
## 8 1841-Harrison  fellow-citizens 11
## 9 1845-Polk      fellow-citizens 1
## 10 1849-Taylor   fellow-citizens 1
## # ... with 45,443 more rows
```

Find the words most specific to each of the inaugural speeches by calculating the tf-idf of each term-speech using the `bind_tf_idf()` function:

```
inaug_tf_idf <- inaug_td %>%
  bind_tf_idf(term, document, count) %>%
  arrange(desc(tf_idf))

inaug_tf_idf
```

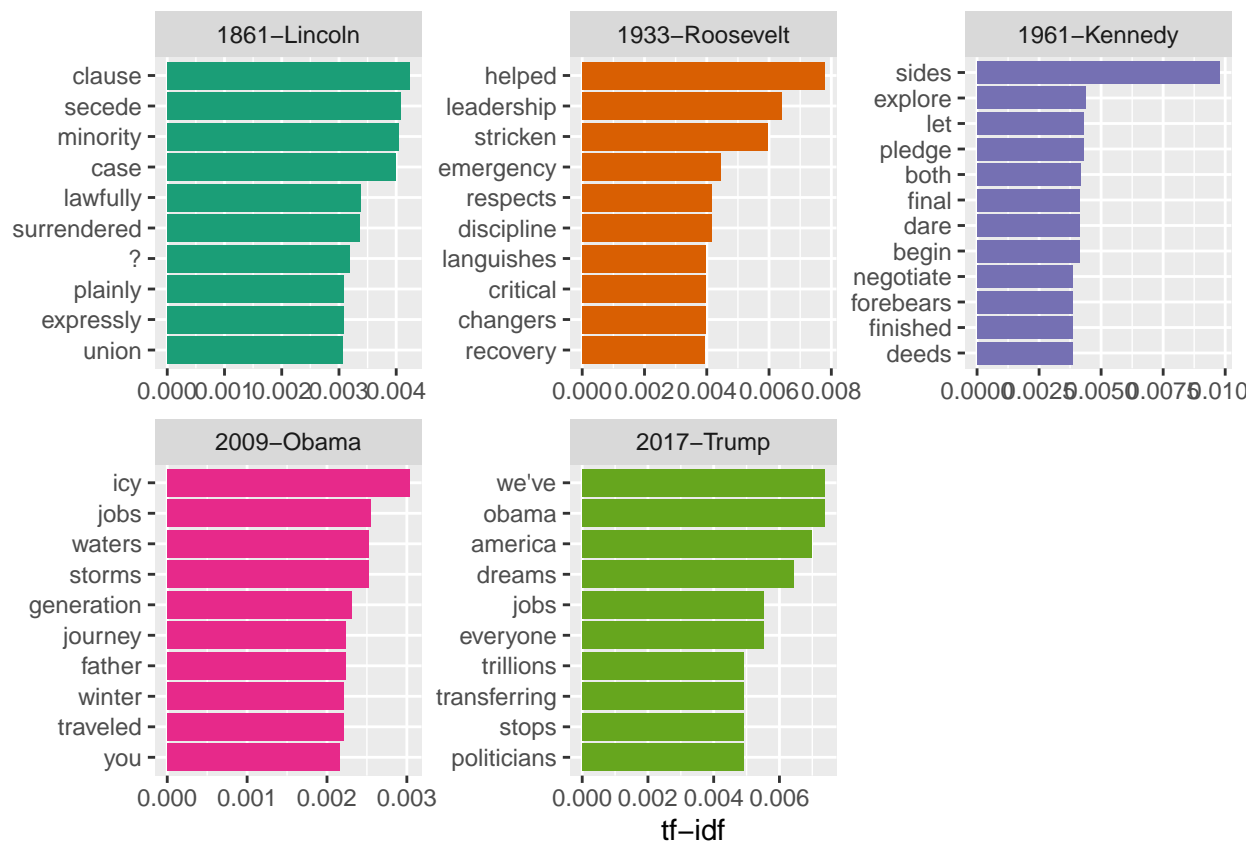
```
## # A tibble: 45,453 x 6
##   document      term      count      tf      idf tf_idf
##   <chr>         <chr>    <dbl>   <dbl> <dbl> <dbl>
## 1 1793-Washington arrive      1 0.00680 4.08 0.0277
## 2 1793-Washington upbraidings 1 0.00680 4.08 0.0277
## 3 1793-Washington violated    1 0.00680 3.38 0.0230
## 4 1793-Washington willingly   1 0.00680 3.38 0.0230
## 5 1793-Washington incurring    1 0.00680 3.38 0.0230
## 6 1793-Washington previous     1 0.00680 2.98 0.0203
## 7 1793-Washington knowingly    1 0.00680 2.98 0.0203
## 8 1793-Washington injunctions  1 0.00680 2.98 0.0203
## 9 1793-Washington witnesses    1 0.00680 2.98 0.0203
## 10 1793-Washington besides     1 0.00680 2.69 0.0183
## # ... with 45,443 more rows
```

```
speeches <- c("1933-Roosevelt", "1861-Lincoln", "1961-Kennedy", "2009-Obama", "2017-Trump")
```

```

inaug_tf_idf %>%
  filter(document %in% speeches) %>%
  group_by(document) %>%
  top_n(10,tf_idf) %>%
  ungroup %>%
  mutate(term=reorder_within(term, tf_idf, document)) %>%
  ggplot(aes(term, tf_idf, fill = document)) +
  geom_col(show.legend = FALSE) +
  facet_wrap(~document, scales = "free") +
  coord_flip() +
  scale_x_reordered() +
  labs(x = NULL,
       y = "tf-idf") +
  scale_fill_brewer(palette = "Dark2")

```



```
# visualize how words changed in frequency over time

year_term_counts <- inaug_td %>%
  extract(document, "year", "(\\d+)", convert = TRUE) %>%
  complete(year, term, fill = list(count = 0)) %>%
  group_by(year) %>%
  mutate(year_total = sum(count))

year_term_counts %>%
  filter(term %in% c("god", "america", "foreign", "union",
                    "trade", "constitution", "freedom", "immigrants",
                    "economy", "education", "environment", "terrorism")) %>%
  ggplot(aes(year, count / year_total)) +
  geom_point() +
  geom_smooth() +
  facet_wrap(~ term, scales = "free_y", ncol = 3) +
  scale_y_continuous(labels = scales::percent_format()) +
  ylab("% frequency of word in inaugural address")
```

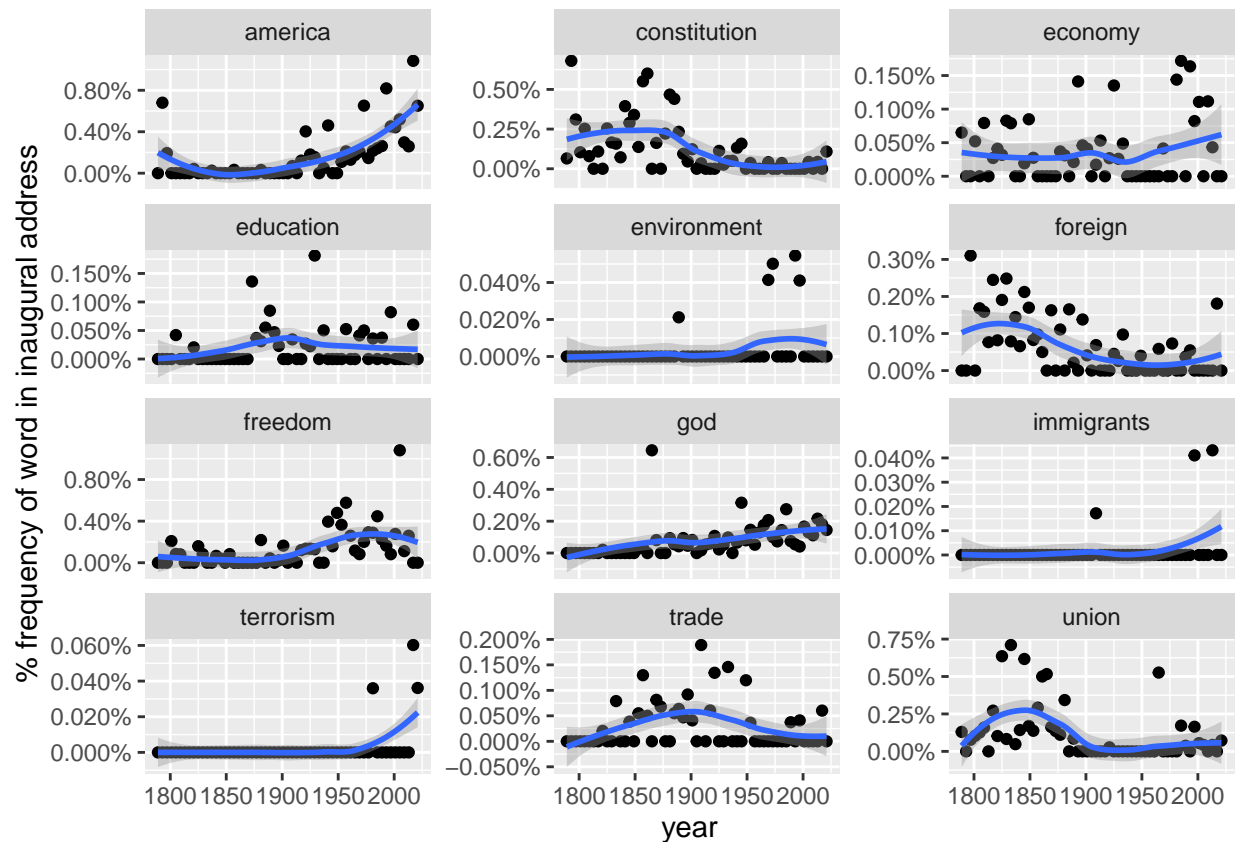


Figure 8: Changes in word frequency over time within Presidential inaugural addresses, for twelve selected terms